

# Lab 2 - Data wrangling

Edmond Niu

```
library(tidyverse)
```

## Part 2

### Question 1

IL has the highest number of counties, at a count of 102. WI has the lowest number of counties, at a count of 72.

```
midwest |>  
  count(state)
```

```
# A tibble: 5 x 2  
  state      n  
  <chr> <int>  
1 IL      102  
2 IN       92  
3 MI       83  
4 OH       88  
5 WI       72
```

## Question 2

Only 3 counties are located in all 5 states in this dataset (Crawford, Jackson, and Monroe).

```
midwest |>
  count(county) |>
  filter(n==5)
```

```
# A tibble: 3 x 2
  county      n
  <chr>    <int>
1 CRAWFORD     5
2 JACKSON      5
3 MONROE       5
```

### Question 3

a.

```
midwest |>
  filter(popdensity > 25000) |>
  select(county, state, popdensity, poptotal, area) |>
  arrange(desc(popdensity))
```

# A tibble: 9 x 5

	county	state	popdensity	poptotal	area
	<chr>	<chr>	<dbl>	<int>	<dbl>
1	COOK	IL	88018.	5105067	0.058
2	MILWAUKEE	WI	63952.	959275	0.015
3	WAYNE	MI	60334.	2111687	0.035
4	CUYAHOGA	OH	54313.	1412140	0.026
5	DU PAGE	IL	39083.	781666	0.02
6	MARION	IN	34659.	797159	0.023
7	HAMILTON	OH	34649.	866228	0.025
8	FRANKLIN	OH	28278.	961437	0.034
9	MACOMB	MI	25621.	717400	0.028

b.

```
midwest |>
  filter(popdensity == max(popdensity)) |>
  select(county, state, popdensity, poptotal, area)
```

# A tibble: 1 x 5

	county	state	popdensity	poptotal	area
	<chr>	<chr>	<dbl>	<int>	<dbl>
1	COOK	IL	88018.	5105067	0.058

## Question 4

The distribution of population density of counties is unimodal and extremely right-skewed. A typical Midwestern county has population density of 1156 people per unit area. The middle 50% of the counties have population densities between 622 to 2330 people per unit area.

```
midwest |>
  summarize(
    median = median(popdensity),
    q1 = quantile(popdensity, 0.25),
    q3 = quantile(popdensity, 0.75)
  )
```

```
# A tibble: 1 x 3
  median    q1    q3
  <dbl> <dbl> <dbl>
1  1156.   622.  2330
```

## Question 5

Proportion of counties in urban areas in each state:

IL: 0.27

IN: 0.40

MI: 0.30

OH: 0.45

WI: 0.28

```
midwest |>
  mutate(metro = if_else(inmetro == 1, "Yes", "No")) |>
  group_by(state, metro) |>
  summarise(count = n()) |>
  mutate(proportion_inmetro = count / sum(count)) |>
  filter(metro == "Yes")
```

`summarise()` has grouped output by 'state'. You can override using the  
`.groups` argument.

```
# A tibble: 5 x 4
# Groups:   state [5]
  state metro count proportion_inmetro
  <chr> <chr> <int>           <dbl>
1 IL    Yes     28           0.275
2 IN    Yes     37           0.402
3 MI    Yes     25           0.301
4 OH    Yes     40           0.455
5 WI    Yes     20           0.278
```

## Question 6

a.

```
midwest |>
  filter(percbelowpoverty == max(percbelowpoverty)) |>
  select(county,state,percbelowpoverty, percollege)
```

# A tibble: 1 x 4

	county	state	percbelowpoverty	percollege
	<chr>	<chr>	<dbl>	<dbl>
1	MENOMINEE	WI	48.7	7.34

b.

```
midwest |>
  filter(percollege > 40) |>
  select(county,state,percbelowpoverty, percollege)
```

# A tibble: 5 x 4

	county	state	percbelowpoverty	percollege
	<chr>	<chr>	<dbl>	<dbl>
1	CHAMPAIGN	IL	15.6	41.3
2	DU PAGE	IL	2.71	42.8
3	HAMILTON	IN	3.59	42.1
4	WASHTENAW	MI	12.2	48.1
5	DANE	WI	10.5	43.6

c.

```
midwest |>
  filter(percollege > 40 | percbelowpoverty == max(percbelowpoverty)) |>
  select(county,state,percbelowpoverty, percollege)
```

# A tibble: 6 x 4

	county	state	percbelowpoverty	percollege
	<chr>	<chr>	<dbl>	<dbl>
1	CHAMPAIGN	IL	15.6	41.3
2	DU PAGE	IL	2.71	42.8

3	HAMILTON	IN	3.59	42.1
4	WASHTENAW	MI	12.2	48.1
5	DANE	WI	10.5	43.6
6	MENOMINEE	WI	48.7	7.34

d.

```
midwest <- midwest |>
  mutate(potential_outlier = if_else(percollege > 40 |
    percbelowpoverty == max(percbelowpoverty), "Yes", "No"))

midwest |>
  select(county, state, percbelowpoverty, percollege, potential_outlier) |>
  arrange(potential_outlier)
```

# A tibble: 437 x 5


	county	state	percbelowpoverty	percollege	potential_outlier
	<chr>	<chr>	<dbl>	<dbl>	<chr>
1	ADAMS	IL	13.2	19.6	No
2	ALEXANDER	IL	32.2	11.2	No
3	BOND	IL	12.1	17.0	No
4	BOONE	IL	7.21	17.3	No
5	BROWN	IL	13.5	14.5	No
6	BUREAU	IL	10.4	18.9	No
7	CALHOUN	IL	15.1	11.9	No
8	CARROLL	IL	11.7	16.2	No
9	CASS	IL	13.9	14.1	No
10	CHRISTIAN	IL	11.7	13.6	No

# i 427 more rows

e.

```
ggplot(midwest, aes(x=percollege, y=percbelowpoverty, color = potential_outlier,
  shape=state)) +
  geom_point() +
  labs(
    x = "percentage of people with a college degree",
    y = "percentage of people below poverty",
    title = "% people with college degree vs people below poverty by state"
  )
```

% people with college degree vs people below poverty by state





## Question 7

a.

```
state_population <- midwest |>
  group_by(state) |>
  summarize(total_population = sum(poptotal))

state_population |>
  arrange(desc(total_population))
```

```
# A tibble: 5 x 2
  state total_population
  <chr>         <int>
1 IL           11430602
2 OH           10847115
3 MI            9295297
4 IN            5544159
5 WI            4891769
```

b.

```
state_population |>
  mutate(propOf_totalPopulation = total_population / sum(total_population)) |>
  arrange(desc(propOf_totalPopulation))
```

```
# A tibble: 5 x 3
  state total_population propOf_totalPopulation
  <chr>         <int>                <dbl>
1 IL           11430602                0.272
2 OH           10847115                0.258
3 MI            9295297                0.221
4 IN            5544159                0.132
5 WI            4891769                0.116
```

c. IL is the most populous Midwestern state with 27.2% of the Midwest population living there. WI is the least populous Midwestern state with 11.6% of the Midwest population living there.

## Question 8

The state that has the lowest average percentage below poverty across its counties is IN (Indiana). The state that has the highest average percentage below poverty across its countries is MI (Michigan).

```
state_poverty <- midwest |>
  group_by(state) |>
  summarize(mean_percbelowpoverty = mean(percbelowpoverty)) |>
  select(state, mean_percbelowpoverty)

state_poverty |>
  arrange(mean_percbelowpoverty)
```

```
# A tibble: 5 x 2
  state mean_percbelowpoverty
  <chr>           <dbl>
1 IN              10.3
2 WI              11.9
3 OH              13.0
4 IL              13.1
5 MI              14.2
```

## Part 2

### Question 9

```
df <- tibble(  
  var_1 = c(10, 20, 30, 40, 50),  
  var_2 = c("Pizza", "Burger", "Pizza", "Pizza", "Burger"),  
  var_3 = c("Apple", "Apple", "Pear", "Pear", "Banana")  
)  
  
df
```

```
# A tibble: 5 x 3  
  var_1 var_2 var_3  
  <dbl> <chr> <chr>  
1     10 Pizza Apple  
2     20 Burger Apple  
3     30 Pizza Pear  
4     40 Pizza Pear  
5     50 Burger Banana
```

a. The code chunk below arranges the values in column `var_2` by alphabetical order. `Arrange()` orders rows using column values in either ascending or descending fashion.

```
df |>  
  arrange(var_2)
```

```
# A tibble: 5 x 3  
  var_1 var_2 var_3  
  <dbl> <chr> <chr>  
1     20 Burger Apple  
2     50 Burger Banana  
3     10 Pizza Apple  
4     30 Pizza Pear  
5     40 Pizza Pear
```

b. The following code groups the data by value in `var_2`. The `group_by()` function groups the data by one or more variables (in this instance, `var_2`). It's different from `arrange` in part a in that `arrange(var_2)` groups the data by column `var_2` but `group_by(var_2)` groups the data by column `var_3` value.

```
df |>
  group_by(var_2)
```

```
# A tibble: 5 x 3
# Groups:   var_2 [2]
  var_1 var_2 var_3
<dbl> <chr> <chr>
1    10 Pizza  Apple
2    20 Burger Apple
3    30 Pizza  Pear
4    40 Pizza  Pear
5    50 Burger Banana
```

c. This code chunk groups the data by value in `var_2` into two groups (Burger and Pizza). It then takes the corresponding `var_1` value and takes the mean of all the `var_1` values from all entries for Burgers and all entries for Pizza and computes two means (for the two groups in `var_2`).

```
df |>
  group_by(var_2) |>
  summarize(mean_var_1 = mean(var_1))
```

```
# A tibble: 2 x 2
  var_2 mean_var_1
<chr>      <dbl>
1 Burger      35
2 Pizza      26.7
```

d. This code chunk groups the data by value in var\_2 AND var\_3 into 4 groups (all combinations of var\_2 values with var\_3 values). It then takes the corresponding var\_1 value(s) from all entries of each combination/grouping of var\_2 and var\_3 values, and computes the mean of the var\_1 value(s).

```
df |>
  group_by(var_2, var_3) |>
  summarize(mean_var_1 = mean(var_1))
```

`summarise()` has grouped output by 'var\_2'. You can override using the `.groups` argument.

```
# A tibble: 4 x 3
# Groups:   var_2 [2]
  var_2 var_3 mean_var_1
  <chr> <chr>      <dbl>
1 Burger Apple        20
2 Burger Banana       50
3 Pizza  Apple        10
4 Pizza  Pear         35
```

e. This code chunk does the exact same thing as part d, but the .groups = “drop” drops all levels of grouping. So this code chunk groups the data by value in var\_2 AND var\_3 into 4 groups (all combinations of var\_2 values with var\_3 values). It then takes the corresponding var\_1 value(s) from all entries of each combination/grouping of var\_2 and var\_3 values, and computes the mean of the var\_1 value(s), BUT at the end, the groupings of var\_2, var\_3 are dropped.

```
df |>
  group_by(var_2, var_3) |>
  summarize(mean_var_1 = mean(var_1), .groups = "drop")
```

```
# A tibble: 4 x 3
  var_2 var_3 mean_var_1
  <chr> <chr>      <dbl>
1 Burger Apple        20
2 Burger Banana       50
3 Pizza  Apple        10
4 Pizza  Pear         35
```

f. The first pipeline does exactly what part e is doing. It groups the data by `var_2` and `var_3` and calculates `var_1` means for each grouping made. It then gets rid of the grouping. However, the second pipeline is different. It also groups by `var_2` and `var_3` but it mutates instead of summarizes. What this means is that each group is NOT summarized down to one row (hence the two rows `mean_var_1` for Pizza and Pear). This is why there are 5 rows instead of 4 for the second pipeline. All `mutate` is doing is changing the values of `mean_var_1` to reflect every single row created by the `group_by(var_2, var_3)`.

```
df |>
  group_by(var_2, var_3) |>
  summarize(mean_var_1 = mean(var_1), .groups = "drop")
```

```
# A tibble: 4 x 3
  var_2 var_3 mean_var_1
  <chr> <chr>      <dbl>
1 Burger Apple      20
2 Burger Banana     50
3 Pizza  Apple      10
4 Pizza  Pear       35
```

```
df |>
  group_by(var_2, var_3) |>
  mutate(mean_var_1 = mean(var_1))
```

```
# A tibble: 5 x 4
# Groups:   var_2, var_3 [4]
  var_1 var_2 var_3 mean_var_1
  <dbl> <chr> <chr>      <dbl>
1    10 Pizza Apple      10
2    20 Burger Apple     20
3    30 Pizza Pear       35
4    40 Pizza Pear       35
5    50 Burger Banana    50
```

### **Question 10**

No answer needed here! Just select questions and pages to indicate where your responses are located when you upload your lab PDF to Gradescope and you'll get full points on this question.